

Report 01.07.20

01/07/2020

Matteo Perotti

Luca Bertaccini

Pasquale Davide Schiavone

Stefan Mach

Professor Luca Benini

Integrated Systems Laboratory

ETH Zürich

Summary

- **OPUS** analysis
- **FP** multiplication

When can HCC solve a size problem?

No initial problem
No big difference

- dijkstra
- fft
- aha-mont64
- minver
- nbody
- nsichneu
- st
- statemate
- wikisort

HCC makes the
difference

- NB-IoT
- sha
- qsort
- crc32
- edn
- matmult-int
- nettle-sha256
- qduino
- sglib-combined
- ud

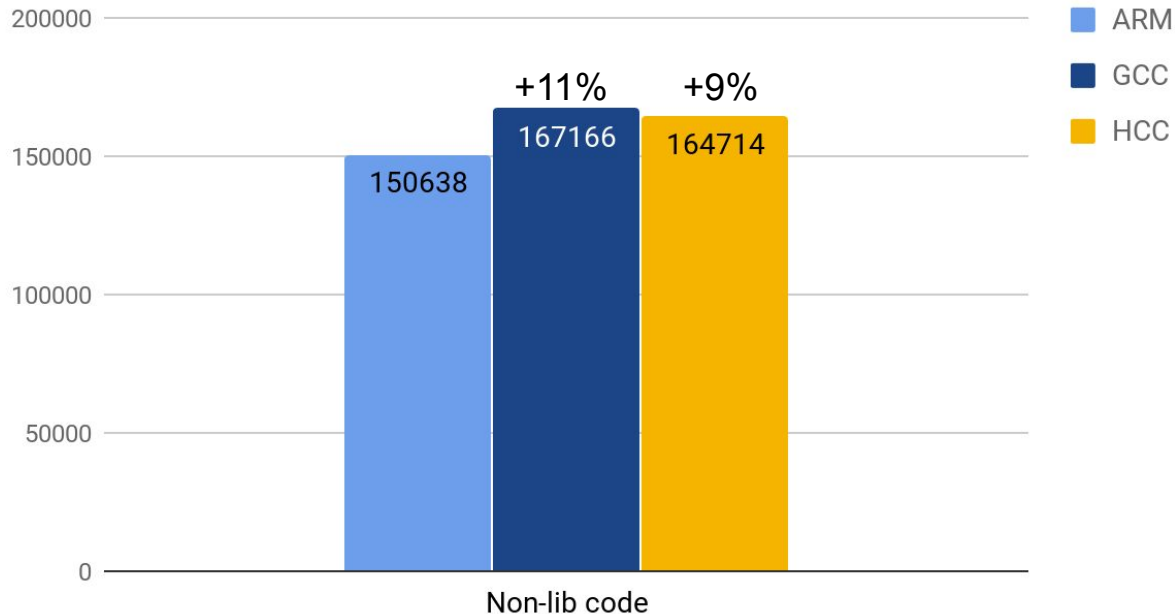
HCC does not solve
the problem

- **opus**
- bitcount
- qsort
- stringsearch
- susan
- huffbench
- nettle-aes
- picojpeg
- slre

opus

- Approximate results (hard to separate all the library functions)
- The size difference is related to non-lib functions
- Analyzed the most inflated functions

opus_demo - code size [B]



opus

OPUS function	ARM size [B]	HCC size [B]	Size diff [B]	Inflation
celt_decode_lost	1796	2986	1190	66%
op_pvq_search_c	660	900	240	36%
validate_celt_decoder	388	596	208	53%

OPUS analysis (<celt_decode_lost>)

- The **function** is **big** and performs a lot of operations
- Hard to precisely make a comparison, but... we can note something!
- **ARM** and **HCC** code have approx the **same number** of **mem ops** (250)
- The **ratio** of **compressed/uncompressed memory ops** is the **same**
- **ARM** always **keeps** the **same register** as **base address** for **+65%** of **memory ops**. This **reg** is **never modified**
- **RISC-V** uses **multiple regs**, every time **modifying** them

OPUS analysis (<celt_decode_lost>)

```
14aca: 6f3b .....ldr r3, [r7, #112] ; 0x70
14acc: 6d7c .....ldr r4, [r7, #84] ; 0x54
14ad0: f107 019c ...add.w r1, r7, #156 ; 0x9c
```

ARM code: 173 mem ops on 256 use r7 as base reg.
r7 is loaded in the beginning and never modified.

```
1dfd8: 777d .....lui a4, 0xffffffff
1dfda: ec870793 .....addi a5, a4, -312
1dfde: 97a2 .....add a5, a5, s0
1dfe0: 438c .....lw a1, 0(a5)
1dfe2: ed470793 .....addi a5, a4, -300
1dfe6: ec070713 .....addi a4, a4, -320
1dfea: 97a2 .....add a5, a5, s0
1dfec: 9722 .....add a4, a4, s0
1dfee: 439c .....lw a5, 0(a5)
1dff0: 4318 .....lw a4, 0(a4)
```

HCC code: before a memory operation, the code loads a value into a reg and adds an offset.

OPUS analysis (<celt_decode_lost>)

- The way **HCC schedules the operations** seems **redundant**.
- In this fragment of code, the **net effect** is to **load** values in **a1, a4, a5**.

```
1dfd8: 777d ..... lui a4,0xffffffff
1dfda: ec870793 ..... addi a5,a4,-312
1dfde: 97a2 ..... add a5,a5,s0
1dfe0: 438c ..... lw a1,0(a5)
1dfe2: ed470793 ..... addi a5,a4,-300
1dfe6: ec070713 ..... addi a4,a4,-320
1dfea: 97a2 ..... add a5,a5,s0
1dfec: 9722 ..... add a4,a4,s0
1dfee: 439c ..... lw a5,0(a5)
1dff0: 4318 ..... lw a4,0(a4)
```

Fragment of <celt_decode_lost> function (compiled with HCC)

OPUS analysis (<celt_decode_lost>)

- The **addresses** are **similar**: we can use only one register as base address

```

1dfd8: 777d ..... lui a4,0xffff
1dfda: ec870793 ..... addi a5,a4,-312
1dfde: 97a2 ..... add a5,a5,s0
1dfe0: 438c ..... lw a1,0(a5)
1dfe2: ed470793 ..... addi a5,a4,-300
1dfe6: ec070713 ..... addi a4,a4,-320
1dfea: 97a2 ..... add a5,a5,s0
1dfec: 9722 ..... add a4,a4,s0
1dfee: 439c ..... lw a5,0(a5)
1dff0: 4318 ..... lw a4,0(a4)

```

Original HCC code

```

1dfd8: 777d ..... lui a4,0xffff
yyyyy: xxxxxxxx ..... addi a4,a4,-320
yyyyy: xxxx ..... add a4,a4,s0
yyyyy: xxxx ..... lw a1,8(a4)
yyyyy: xxxx ..... lw a5,20(a4)
yyyyy: xxxx ..... lw a4,0(a4)

```

Manually optimized code

OPUS analysis (<celt_decode_lost>)

- Another portion of the code (GCC 9.2 does not solve the issue):

```

1dffc: 76fd.....lui a3,0xfffff
1dffe: ed468793.....addi a5,a3,-300.#
1e002: ec068713.....addi a4,a3,-320
1e006: 97a2.....add a5,a5,s0
1e008: 9722.....add a4,a4,s0
1e00a: 4318.....lw a4,0(a4)
1e00c: 439c.....lw a5,0(a5)
1e00e: 89b6.....mv s3,a3
1e010: e9c98613.....addi a2,s3,-356
1e014: 60f717db.....mulhiadd a5,a4,a5,96
1e018: eb068713.....addi a4,a3,-336
1e01c: 9722.....add a4,a4,s0
1e01e: ed868693.....addi a3,a3,-296
1e022: eb098593.....addi a1,s3,-336
1e026: eb498513.....addi a0,s3,-332
1e02a: c31c.....sw a5,0(a4)
1e02c: 96a2.....add a3,a3,s0
1e02e: 9622.....add a2,a2,s0
1e030: 95a2.....add a1,a1,s0
1e032: 9522.....add a0,a0,s0
1e034: 50dc.....lw a5,36(s1)
1e036: 4294.....lw a3,0(a3)
1e038: 4210.....lw a2,0(a2)
1e03a: 418c.....lw a1,0(a1)
1e03c: 4108.....lw a0,0(a0)
1e03e: 4761.....li a4,24

```

Original HCC code

Manual reordering



```

1dffc: 76fd.....lui a3,0xfffff
1dffe: ed468793.....addi a5,a3,-300.#
1e002: ec068713.....addi a4,a3,-320
1e006: 97a2.....add a5,a5,s0
1e008: 9722.....add a4,a4,s0
1e00a: 4318.....lw a4,0(a4)
1e00c: 439c.....lw a5,0(a5)
1e014: 60f717db.....mulhiadd a5,a4,a5,96
1e00e: 89b6.....mv s3,a3
1e018: eb068713.....addi a4,s3,-336
1e01e: ed868693.....addi a3,s3,-296
1e010: e9c98613.....addi a2,s3,-356
1e022: eb098593.....addi a1,s3,-336
1e026: eb498513.....addi a0,s3,-332
1e01c: 9722.....add a4,a4,s0
1e02c: 96a2.....add a3,a3,s0
1e02e: 9622.....add a2,a2,s0
1e030: 95a2.....add a1,a1,s0
1e032: 9522.....add a0,a0,s0
1e02a: c31c.....sw a5,0(a4)
1e034: 50dc.....lw a5,36(s1)
1e036: 4294.....lw a3,0(a3)
1e038: 4210.....lw a2,0(a2)
1e03a: 418c.....lw a1,0(a1)
1e03c: 4108.....lw a0,0(a0)
1e03e: 4761.....li a4,24

```

Equivalent code

OPUS analysis (<celt_decode_lost>)

```

1dff: 76fd ..... lui a3,0xffff
1dfe: ed468793 ..... addi a5,a3,-300 #
1e00: ec068713 ..... addi a4,a3,-320
1e06: 97a2 ..... add a5,a5,s0
1e08: 9722 ..... add a4,a4,s0
1e0a: 4318 ..... lw a4,0(a4)
1e0c: 439c ..... lw a5,0(a5)
1e14: 60f717db ..... muliadd a5,a4,a5,96
1e0e: 89b6 ..... mv s3,a3
1e18: eb068713 ..... addi a4,s3,-336
1e1e: ed868693 ..... addi a3,s3,-296
1e10: e9c98613 ..... addi a2,s3,-356
1e22: eb098593 ..... addi a1,s3,-336
1e26: eb498513 ..... addi a0,s3,-332
1e1c: 9722 ..... add a4,a4,s0
1e2c: 96a2 ..... add a3,a3,s0
1e2e: 9622 ..... add a2,a2,s0
1e30: 95a2 ..... add a1,a1,s0
1e32: 9522 ..... add a0,a0,s0
1e2a: c31c ..... sw a5,0(a4)
1e34: 50dc ..... lw a5,36(s1)
1e36: 4294 ..... lw a3,0(a3)
1e38: 4210 ..... lw a2,0(a2)
1e3a: 418c ..... lw a1,0(a1)
1e3c: 4108 ..... lw a0,0(a0)
1e3e: 4761 ..... li a4,24

```

Equivalent code

- The memory operations are:
 - `lw rd, 0(0xffff000 - off + s0)`
- **off** is:
 - 296
 - 300
 - 320
 - 332
 - 336
 - 336
 - 356
- **off** is always in the range [296, 356]
- The **sequence** of “addi + add” is **repeated** for **each register** (a0, a1, a2, a3, a4)
- But with the **load/store** operations we can **add** an **offset** **in place!!!**

OPUS analysis (<celt_decode_lost>:)

- Strange compiler choices (GCC 9.2 produces similar code)

```

1dfffc: 76fd.....lui a3,0xfffff
1dfffe: ed468793.....addi a5,a3,-300 #
1e002: ec068713.....addi a4,a3,-320
1e006: 97a2.....add a5,a5,s0
1e008: 9722.....add a4,a4,s0
1e00a: 4318.....lw a4,0(a4)
1e00c: 439c.....lw a5,0(a5)
1e014: 60f717db.....muladd a5,a4,a5,96
1e01e: 89b6.....mv s3,a3
1e018: eb068713.....addi a4,s3,-336
1e01e: ed868693.....addi a3,s3,-296
1e010: e9c98613.....addi a2,s3,-356
1e022: eb098593.....addi a1,s3,-336
1e026: eb498513.....addi a0,s3,-332
1e01c: 9722.....add a4,a4,s0
1e02c: 96a2.....add a3,a3,s0
1e02e: 9622.....add a2,a2,s0
1e030: 95a2.....add a1,a1,s0
1e032: 9522.....add a0,a0,s0
1e02a: c31c.....sw a5,0(a4)
1e034: 50dc.....lw a5,36(s1)
1e036: 4294.....lw a3,0(a3)
1e038: 4210.....lw a2,0(a2)
1e03a: 418c.....lw a1,0(a1)
1e03c: 4108.....lw a0,0(a0)
1e03e: 4761.....li a4,24

```

Equivalent code

- ARM: 40 B
- HCC: **68 B**
- Manually optimized: **30 B**

Manual optimization



```

1dfffc: 76fd.....lui a0,0xfffff
yyyyy: xxxx.....add a0,a0,s0
yyyyy: xxxxxxxx.....addi a0,a0,-356
1e00c: 439c.....lw a5,56(a0)
1e00a: 4318.....lw a4,36(a0)
1e014: 60f717db.....muladd a5,a4,a5,96
1e02a: c31c.....sw a5,20(a0)
1e034: 50dc.....lw a5,36(s1)
1e036: 4294.....lw a3,60(a0)
1e038: 4210.....lw a2,0(a0)
1e03a: 418c.....lw a1,20(a0)
1e03c: 4108.....lw a0,24(a0)
1e03e: 4761.....li a4,24

```

Equivalent code

OPUS analysis (<celt_decode_lost>:)

```

1dfd8: 777d ..... lui a4,0xffff
yyyyy: xxxxxxxx ..... addi a4, a4, -320
yyyyy: xxxx ..... add a4, a4, s0
yyyyy: xxxx ..... lw a1, 8(a4)
yyyyy: xxxx ..... lw a5, 20(a4)
yyyyy: xxxx ..... lw a4, 0(a4)

```

```

1dffc: 76fd ..... lui a0,0xffff
yyyyy: xxxx ..... add a0, a0, s0
yyyyy: xxxxxxxx ..... addi a0, a0, -356
1e00c: 439c ..... lw a5, 56(a0)
1e00a: 4318 ..... lw a4, 36(a0)
1e014: 60f717db ..... muliadd a5, a4, a5, 96
1e02a: c31c ..... sw a5, 20(a0)
1e034: 50dc ..... lw a5, 36(s1)
1e036: 4294 ..... lw a3, 60(a0)
1e038: 4210 ..... lw a2, 0(a0)
1e03a: 418c ..... lw a1, 20(a0)
1e03c: 4108 ..... lw a0, 24(a0)
1e03e: 4761 ..... li a4, 24

```

- These are the **two** manually **optimized fragments** of code.
- $s0 = sp + \text{fixed_offset}$
- After the initial assignment, **s0** is **never modified**.
- We always **load 0xfffff**, then **add s0** and a **variable offset**.
- This **pattern** is **frequent** into the function (it seems >60 times).
- The **variable offsets** are very **similar** also for **different fragments**.



- A **fixed register** for the **whole function** as a **base** for the memory operations (as ARM does) can be **convenient**.

OPUS analysis (<op_pvq_search_c>)

- >50 memory operations with **s0** as a **base register**
- **s0** is a **fixed value** loaded in the prologue of the function (sp + offset)
- This fixed value **imposes** a **negative offset** in the memory ops
- The negative offset **prevents** the **compression** of the **memory ops**
- Tried with **GCC 9.2**, this problem seems **solved!**

```

23e74:-fac42823.....sw—a2,-80(s0)
23e78:-fad42a23.....sw—a3,-76(s0)
23e7c:-fae42c23.....sw—a4,-72(s0)
23e80:-faf42e23.....sw—a5,-68(s0)
23e84:-766180ef.....jal—ra,3c5ea<
23e88:-fb842703.....lw—a4,-72(s0)
23e8c:-fac42803.....lw—a6,-84(s0)
23e90:-00000793.....li—a5,0
23e94:-fb042603.....lw—a2,-80(s0)
23e98:-fb442683.....lw—a3,-76(s0)

```

Fragment of HCC code (<op_pvq_search_c>)

OPUS analysis (<validate_celt_decoder>)

```
00014f90:<validate_celt_decoder>:
```

```
·14fb8:-2277·····—movs—r2,·#119—;·0x77
→·14fba:-4941·····—ldr—r1,·[pc,·#260]—;·(150c0:
·14fbc:-4842·····—ldr—r0,·[pc,·#264]—;·(150c8:
·14fbe:-e7f6·····—b.n-14fae·<validate_celt_deco
```

```
·150ba:-e778·····—b.n-14fae·<validate_celt
·150bc:-bd38·····—pop—{r3,·r4,·r5,·pc}
·150be:-bf00·····—nop
(150c0:-000358f6)·—strdeq—r5,·[r3],·-r6
·150c4:-0003593a·—andeq—r5,·r3,·sl,·lsr·r9
·150c8:-00035982·—andeq—r5,·r3,·r2,·lsr·#19
·150cc:-000359a7·—andeq—r5,·r3,·r7,·lsr·#19
·150d0:-0003338c·—andeq—r3,·r3,·ip,·lsr·#7
```

This value is loaded 20 times

- Alternative to 48-bit l.li: **constant pool** after the function code
- Byte cost per mem op, on N mem ops:
 - PC load*: $2\text{ B} + (4/N)\text{ B} \approx 2\text{ B}$ if N is high
 - l.li*: **6 B**
- Implementation cost: **PC relative loads** (almost for free) vs. **48-bit support** (mod. aligner + ID)
- Drawback: increased **D-memory traffic** (but high locality)

Further

- Go on with the **analysis** of **OPUS**
- Finish the **FP multiplication** (almost done)
- Start to code the **FMA** for HW comparison

Tiny Floating-Point Unit

Tiny FPU:

- Final **FMA** design
- Added **COMPARISON** instructions + optimization
- Added **CAST** instructions (to be optimized)

FMA HW optimizations - Single Precision

FP32	fpnew_fma (NON-REDUCED)	floatli_fma (no_input regs)	floatli_fma_input_regs	floatli_fma_input_regs_ reused
Overall Area	100%	51.2%	60.9%	58.2%
Comb Area	100%	38.8%	41.0%	44.2%
Non-comb	0%	12.4%	19.9%	14.0%
Latency	1 cycle (ADD/MUL/FMADD)	9-11 cycles (ADD) 21-23 cycles (FMADD/MUL)	9-11 cycles (ADD) 21-23 cycles (FMADD/MUL)	9-11 cycles (ADD) 21-23 cycles (FMADD/MUL)
Optimization	-	~48.8%	~39.1%	~41.8%


FMA HW optimizations - Double Precision

FP64	fpnew_fma (NON-REDUCED)	floatli_fma (no_input regs)	floatli_fma_input_regs	floatli_fma_input_regs_ reused
Overall Area	100%	33.6%	39.8%	37.9%
Comb Area	100%	25.4%	27.2%	29.5%
Non-comb	0%	8.2%	12.6%	8.4%
Latency	1 cycle (ADD/MUL/FMADD)	9-11 cycles (ADD) 35-37 cycles (FMADD/MUL)	9-11 cycles (ADD) 35-37 cycles (FMADD/MUL)	9-11 cycles (ADD) 35-37 cycles (FMADD/MUL)
Optimization	-	~66.4%	~60.2%	~62.1%

Non-reduced Floating-Point Unit modules

- **FMA** (FMADD, FNMSUB, ADD, SUB, MUL)
- **COMPARISON** (SGNJ, MINMAX, CMP, CLASSIFY)
- **CAST** :
 - Single-precision: $\text{FP32} \longleftrightarrow \text{INT32}$
 - Double-precision:
 - $\text{FP64} \longleftrightarrow \text{INT64}$
 - $\text{FP64} \longleftrightarrow \text{INT32}$
 - $\text{FP32} \longleftrightarrow \text{INT64}$
 - $\text{FP32} \longleftrightarrow \text{INT32}$

Non-reduced Floating-Point Unit



FMA + COMP HW optimizations - Single Precision

FP32	fpnew_fma + fpnew_noncomp (NON-REDUCED)	reduced_fma_comp	reduced_fma_reduced _comp	reduced_fma_reduced _comp_idle
Overall Area	100%	59.7%	58.2%	53.7%
Latency	1 cycle (ADD/MUL/FMADD/ COMP)	1 cycle (COMP) 9-11 cycles (ADD) 21-23 cycles (FMADD/MUL)	1 cycle COMP 9-11 cycles (ADD) 21-23 cycles (FMADD/MUL)	2 cycle COMP 10-12 cycles (ADD) 22-24 cycles (FMADD/MUL)
Optimization	-	~40.3%	~41.8%	~46.3%


- No additional non-combinational area

FMA + COMP HW optimizations - Double Precision

FP64	fpnew_fma + fpnew_noncomp (NON-REDUCED)	reduced_fma_comp	reduced_fma_reduced _comp	reduced_fma_reduced _comp_idle
Overall Area	100%	39.2%	38.4%	35.9%
Latency	1 cycle (ADD/MUL/FMADD/ COMP)	1 cycle (COMP) 9-11 cycles (ADD) 35-37 cycles (FMADD/MUL)	1 cycle COMP 9-11 cycles (ADD) 35-37 cycles (FMADD/MUL)	2 cycle COMP 10-12 cycles (ADD) 36-38 cycles (FMADD/MUL)
Optimization	-	~60.8%	~61.6%	~64.1%

- No additional non-combinational area

FMA + COMP HW optimizations



FMA + COMP + CAST HW optimizations - Single Precision

FP32	fpnew_fma + fpnew_noncomp + fpnew_multi_cast (NON-REDUCED)	reduced fma_reduced comp_cast_idle
Overall Area	100%	~59.5%
Latency	1 cycle (ADD/MUL/FMADD/COMP)	2 cycle COMP/CAST 10-12 cycles (ADD) 22-24 cycles (FMADD/MUL)
Optimization	-	~40.5%

FMA + COMP + CAST HW optimizations - Double Precision

FP64	fpnew_fma + fpnew_noncomp + fpnew_multi_cast (NON-REDUCED)	reduced fma_reduced comp_cast_idle
Overall Area	100%	~43.0%
Latency	1 cycle (ADD/MUL/FMADD/COMP)	2 cycle COMP/CAST 10-12 cycles (ADD) 36-38 cycles (FMADD/MUL)
Optimization	-	~57.0%

Next steps

- Re-use of FMA components for multicycle CAST:
 - Shifter
 - Adder
 - Registers