

Report 12.08.20

12/08/2020

Matteo Perotti

Luca Bertaccini

Pasquale Davide Schiavone

Stefan Mach

Professor Luca Benini

Integrated Systems Laboratory

ETH Zürich

Tiny Floating-Point Unit

- **Architecture: SNITCH**
- **Tiny FPU: Top-level**
- **Microcode**
- **Target Conference: DATE21**

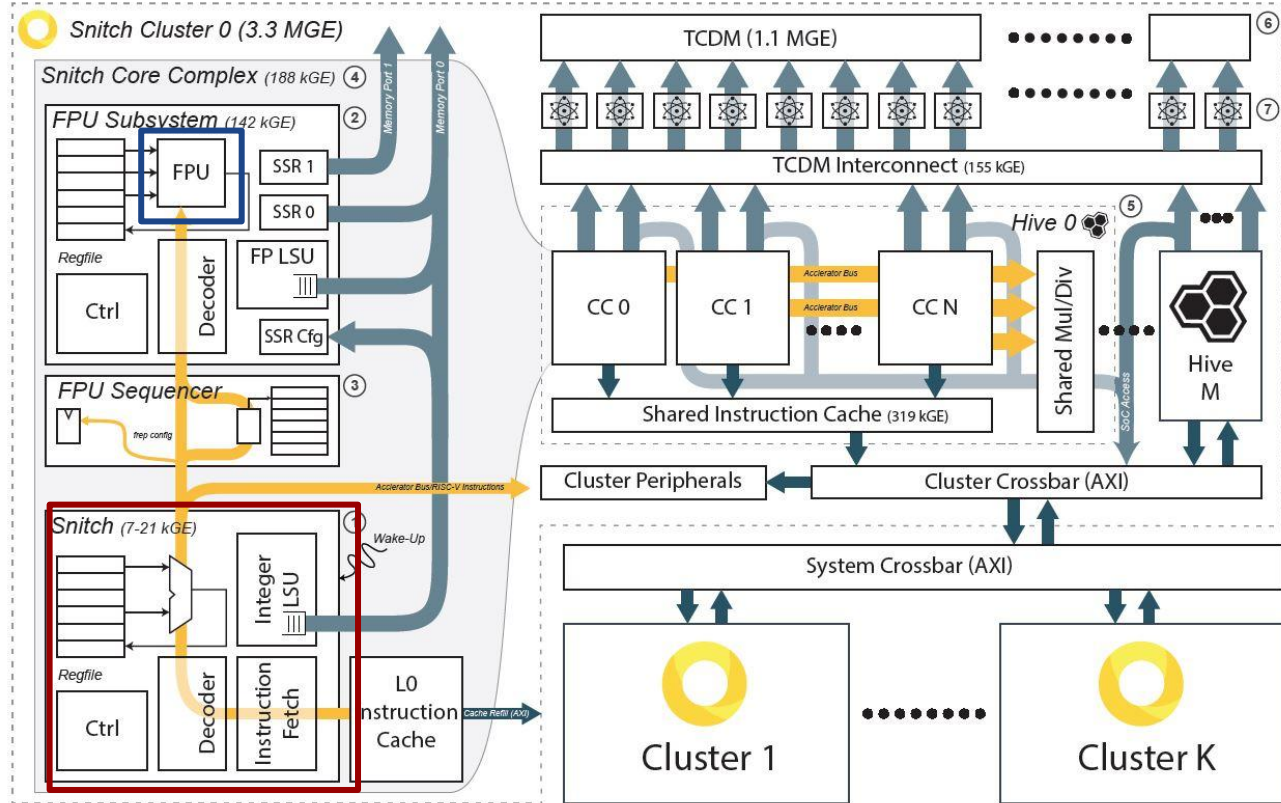
SNITCH

- **Snitch RISC-V Processing Element with Streaming Registers (RV32[I|E]MAFDZifenceiZicsrXssrXfrep)**
 - Support for FP32 and FP64
 - **32-bit INT register file**
 - **64-bit FP register file**
- Snitch paper: <https://arxiv.org/abs/2002.10143>
- *“A general-purpose, **single-stage, single-issue core** tuned for utmost **energy efficiency**. Aiming to maximize the compute/control ratio (making the **FPU the dominant part of the design**) mitigating the effects of deep pipelines and dynamic scheduling”*

SNITCH

Tiny FPU

Snitch

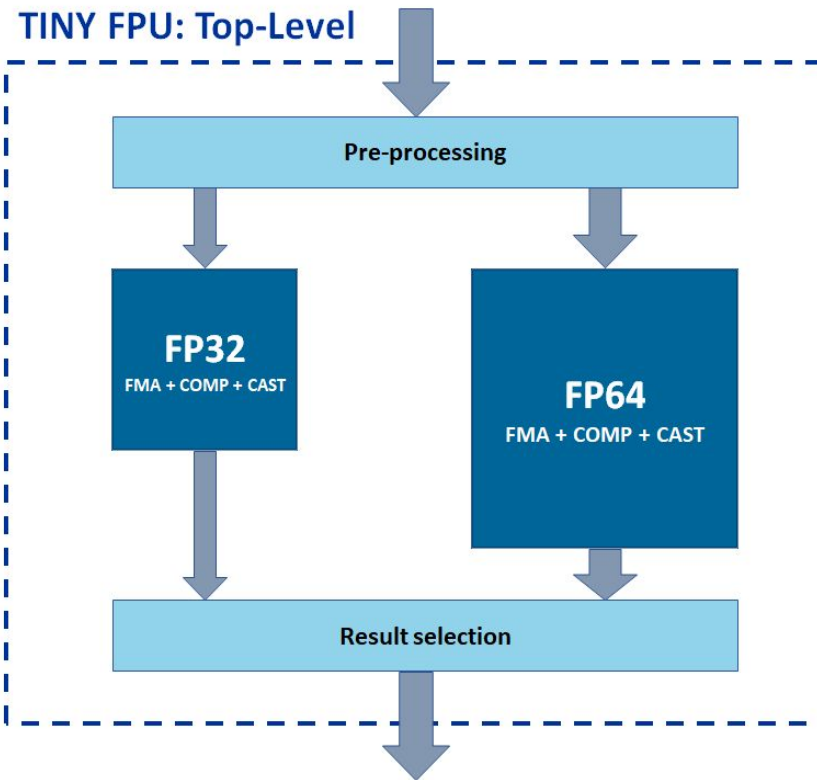


Tiny FPU: Top-Level

- Same **interface** as **FPNEW**
- The **top-level** contains:
 - 1 instantiation of **FP32** FMA+COMP+CAST
 - 1 instantiation of **FP64** FMA+COMP+CAST
 - **pre/post-processing** depending on the instruction and on the precision
- Tiny FPU **integrated into Snitch**

Tiny FPU: Top-Level

- **FP64 module** instantiated just if **double-precision** is **supported**
- **Pre-processing**: depending on the instruction some operands may be set to zero
- **Result selected** depending on input and output **format**



Microcode

- Compile with **F/D extensions** enabled (e.g. fadd.s f1,f2,f3)
- Employ a memory-mapped **ROM** to store the SW emulation instructions, which uses RV32 instructions
- FP instruction will be interpreted by the decoder as a ***jal*** to the ROM address

Microcode

Issue:

- RISC-V ISA does not contain MV from FP64 register file to INT32 register file

Proposed solution:

- Add two custom instructions to perform the move from FP64 to INT32 and vice versa
 - Push & Pop of the integer registers used by the microcode
- or
- Add some 64-bit **shadow registers** to save function call overhead

Target Conference: DATE21

DATE21:

- **Abstract** submission: **September, 14**
- **Paper** submission: **September, 21**

DATE21: Paper content

- Comparison among:
 - **FPU**
 - **Tiny FPU**
 - **SW emulation**
 - **Microcode**
- Performance:
 - Latency
 - Area
 - Code size
 - Power consumption

Next steps

- Microcode implementation
- Optimized SW libraries
- Benchmarks (fft, matmul, knn, ...)

Appendix: FMA + COMP + CAST HW - Single Precision

FP32	fpnew_fma + fpnew_noncomp + fpnew_multi_cast (NON-REDUCED)	reduced_fma_comp_cast (last presentation)	reduced_fma_comp_cast tested
Overall Area	100%	~57.1%	~58.3%
Latency	1 cycle (ADD/MUL/FMADD/COMP)	2 cycles COMP 9 cycles CAST 10-12 cycles (ADD) 22-24 cycles (FMADD/MUL)	2 cycles COMP 9 cycles CAST 10-12 cycles (ADD) 22-24 cycles (FMADD/MUL)
Optimization	-	~42.9%	~41.7%


- No changes in the latencies after testing

Appendix: FMA + COMP + CAST HW - Double Precision

FP64	fpnew_fma + fpnew_noncomp + fpnew_multi_cast (NON-REDUCED)	reduced_fma_comp_cast (last presentation)	reduced_fma_comp_cast tested
Overall Area	100%	~39.3%	~39.6%
Latency	1 cycle (ADD/MUL/FMADD/COMP)	2 cycles COMP 9 cycles CAST 10-12 cycles (ADD) 36-38 cycles (FMADD/MUL)	2 cycles COMP 9 cycles CAST 10-12 cycles (ADD) 36-38 cycles (FMADD/MUL)
Optimization	-	~60.7%	~60.4%

- No changes in the latencies after testing

Appendix: FMA + COMP + CAST: HW optimization (tested)



Appendix: FMA - Single Precision

FP32	Latency reduced FPU	Average Latency SW emulation (SEGGER)	Code size (SEGGER)	Speed-up
ADD	10-12 cycles	49.5 cycles	410 B	~4.5x
SUB	10-12 cycles	62.2 cycles	10 B	~5.5x
MUL	22-24 cycles	39.3 cycles	178 B	~1.7x
FMADD	22-24 cycles	$49.5 + 39.3 = 88.8$ cycles		~3.9x

<https://www.segger.com/products/development-tools/runtime-library/technology/floating-point-library/>

Appendix: COMP - Single Precision

FP32	Latency reduced FPU	Average Latency SW emulation (SEGGER)	Code size (SEGGER)	Speed-up
<	2 cycles	11 cycles	58 B	5.5x
<=	2 cycles	10 cycles	54 B	5x
>	2 cycles	10 cycles	50 B	5x
>=	2 cycles	11 cycles	62 B	5.5x
==	2 cycles	10 cycles	44 B	5x
!=	2 cycles	10 cycles	--	5x

<https://www.segger.com/products/development-tools/runtime-library/technology/floating-point-library/>

Appendix: CAST - Single Precision

FP32	Latency reduced FPU	Average Latency SW emulation (SEGGER)	Code size (SEGGER)	Speed-up
INT32 -> FP32	9 cycles	32.6 cycles	66 B	3.6x
FP32 -> INT32	9 cycles	14 cycles	74 B	1.5x

Overall code size - SW emulation (FMA+COMP+CAST) = **1006 B**

<https://www.segger.com/products/development-tools/runtime-library/technology/floating-point-library/>

Appendix: FMA - Double Precision

FP64	Latency reduced FPU	Average Latency SW emulation (SEGGER)	Code size (SEGGER)	Speed-up
ADD	10-12 cycles	62.8 cycles	724 B	~5.7x
SUB	10-12 cycles	82.8 cycles	10 B	~7.5x
MUL	36-38 cycles	75.0 cycles	286 B	~2x
FMADD	36-38 cycles	62.8 + 75 = 137.8 cycles	-	~3.7x

Overall code size - SW emulation (FMA+COMP+CAST) = **2332 B**

<https://www.segger.com/products/development-tools/runtime-library/technology/floating-point-library/>

Appendix: COMP - Double Precision

FP64	Latency reduced FPU	Average Latency SW emulation (SEGGER)	Code size (SEGGER)	Speed-up
<	2 cycles	16 cycles	70 B	~8x
<=	2 cycles	16 cycles	70 B	~8x
>	2 cycles	16.1 cycles	70 B	~8x
>=	2 cycles	16.1 cycles	70 B	~8x
==	2 cycles	14 cycles	52 B	~7x
!=	2 cycles	14 cycles	--	~7x

<https://www.segger.com/products/development-tools/runtime-library/technology/floating-point-library/>

Appendix: CAST - Double Precision

FP64	Latency reduced FPU	Average Latency SW emulation (SEGGER)	Code size (SEGGER)	Speed-up
INT32 -> FP32	9 cycles	32.6 cycles	66 B	~3.6x
FP32 -> INT32	9 cycles	14 cycles	74 B	~1.5x
INT64 -> FP32	9 cycles	49.1 cycles	96 B	~5.5x
FP32 -> INT64	9 cycles	23.2 cycles	146 B	~2.6x
FP32 -> FP64	9 cycles	14.1 cycles	64 B	~1.5x
FP64 -> INT64	9 cycles	26.9 cycles	146 B	~3.0x
FP64 -> INT32	9 cycles	16.8 cycles	84 B	~3.6x
INT32 -> FP64	9 cycles	31.6 cycles	46 B	~1.9x
INT64 -> FP64	9 cycles	45.1 cycles	128 B	~5x
FP64 -> FP32	9 cycles	25.1 cycles	130 B	~2.8x