

Report 15.10.20

15/10/2020

Matteo Perotti

Luca Bertaccini

Pasquale Davide Schiavone

Stefan Mach

Professor Luca Benini

Integrated Systems Laboratory

ETH Zürich

Summary

- Completed **testing** environment for **NoDenormal** functions
- Further **optimized code** after performance evaluation
- **Completing** the **library** (division, cast functions)
- Environment to **benchmark Tiny FPU** and evaluate its **energy efficiency**

Testing environment for NoDenormal

- Completed testing environment:

- **RV32IM + Denormals**



Verified with TestFloat
(level 2, max number of
inputs)

- RV32IM No Denormals



- RV32EM No Denormals

Verified with TestFloat
(level 2, max number of
inputs)

Updated Code size

| Code Size (B) | Mine | libgcc | Mine ND (rv32em) | SEGGER ND * |
|---------------|-------------|--------|------------------|-------------|
| f32_add | 414 | 786 | 334 | 410 |
| f32_sub | 8 | 816 | 8 | 10 |
| f32_mul | 374 | 538 | 172 | 178 |
| f32_eq | 38 | 86 | 38 | 44 |
| f32_le | 60 | 120 | 60 | 54 |
| f32_lt | 52 | 120 | 52 | 58 |
| f64_add | 750 | 1520 | 572 | 724 |
| f64_sub | 10 | 1538 | 10 | 10 |
| f64_mul | 560 | 1096 | 288 | 286 |
| f64_eq | 64 | 106 | 46 | 52 |
| f64_le | 104 | 166 | 84 | 70 |
| f64_lt | 92 | 166 | 76 | 70 |
| Total | 2526 (-64%) | 7058 | 1740 (-11.5%) | 1966 |

Updated Code size

| Code Size (B) | Mine | libgcc | Mine ND (rv32em) | SEGGER ND * |
|---------------|--------------------|-------------|----------------------|-------------|
| f32_add | 414 | 786 | 334 | 410 |
| f32_sub | 8 | 816 | 8 | 10 |
| f32_mul | 374 | 538 | 172 | 178 |
| f32_eq | 38 | 86 | 38 | 44 |
| f32_le | 60 | 120 | 60 | 54 |
| f32_lt | 52 | 120 | 52 | 58 |
| f64_add | 750 | 1520 | 572 | 724 |
| f64_sub | 10 | 1538 | 10 | 10 |
| f64_mul | 560 | 1096 | 288 | 286 |
| f64_eq | 64 | 106 | 46 | 52 |
| f64_le | 104 | 166 | 84 | 70 |
| f64_lt | 92 | 166 | 76 | 70 |
| Total | 2526 (-64%) | 7058 | 1740 (-11.5%) | 1966 |

Updated performance

| | Mine (cycles) | libgcc (cycles) | Improvement |
|---------|---------------|-----------------|-------------|
| f32_add | 75.20 | 95.50 | 26.99% |
| f32_sub | 79.00 | 102.80 | 30.13% |
| f32_mul | 51.20 | 117.20 | 128.91% |
| f32_eq | 15.00 | 27.30 | 82.00% |
| f32_le | 17.70 | 36.80 | 107.91% |
| f32_lt | 16.70 | 36.80 | 120.36% |
| f64_add | 102.40 | 115.00 | 12.30% |
| f64_sub | 107.10 | 123.50 | 15.31% |
| f64_mul | 91.40 | 219.50 | 140.15% |
| f64_eq | 24.00 | 31.30 | 30.42% |
| f64_le | 25.00 | 40.10 | 60.40% |
| f64_lt | 25.00 | 40.10 | 60.40% |

Notes

- Mine vs. SEGGER:
 - Code size improvement (-11.5%)
 - Lower code size for the addition
 - Lower performance for the addition
- Mine vs. libgcc:
 - Code size improvement (-64%)
 - The real improvement is higher, as libgcc depends on count leading zero function and table!
 - Multiplication/comparison performance boost (also +140% speed)

Completing the library

Functions to add:

- Division (32b-64b)
- Conversion (32b-64b)
 - Float <--> int
 - Float <--> unsigned integer
 - Float <--> signed long
 - Float <--> unsigned long

Open-Source Release

Future plans:

- **Release** the **library** under **GPL** license
- **Complete** and **test** the missing functions
- **Add** them to the library

Benchmark environment

Comparison between FastFPU, TinyFPU, uCode, SW-libgcc:

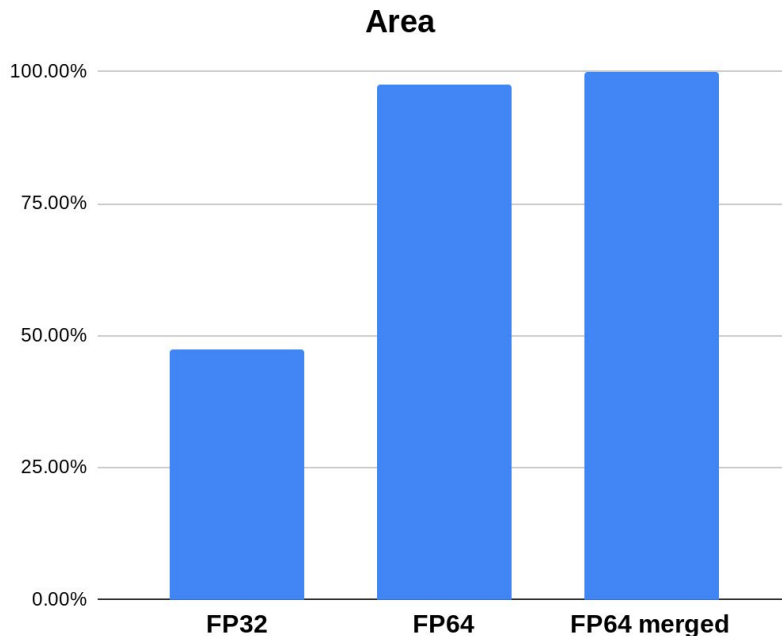
- Estimate **Power, FP-Throughput, Energy Efficiency** (y-axis)
- Vary **% FP operations** in the program (x-axis)
- 1 **synthetic** and some **real** programs from CoreMarkPRO
(<https://github.com/eembc/coremark-pro/tree/master/benchmarks/fp>)
- Find **intervals** of **% FP** operations in which **one implementation** is **more energy efficient** than the others

Tiny Floating-Point Unit

- Merged implementation
- FPU power analysis
- Architectures


TinyFPU - Merged Implementation (Area Overhead)

- Target: **RVFD**
- Merged implementation instead of FP32 TinyFPU and FP64 TinyFPU




TinyFPU - Merged Implementation (Area Overhead)

- Target: **RVFD**
- Merged implementation instead of FP32 TinyFPU and FP64 TinyFPU
- **3% area overhead**
- **+1 cycle** (FMADD, FMUL)



TinyFPU - Merged Implementation



- **FP32 exp** mapped into **FP64 exp** LSBs
- **FP32 mantissa** mapped into **FP64 mantissa** MSBs
- FP32 mantissa + hidden one = 24 bits
- FP64 mantissa + hidden one = 53 bits

TinyFPU vs FastFPU - Power analysis

- Topographical synthesis
- Stand-alone FPUs: power analysis (relaxed timing - **Tclk=15000ps**):
 - **LKG** power **scales** similarly to **area**
 - **Overall power scales** similarly to **area** (FP32 vs FP64 version of the **same FPU**)
 - FP64 TinyFPU is **2.2x** smaller than FP64 FastFPU but consumes **5.5x** less power
 - Lower **switching activity** for TinyFPU

Snitch

- **IPU, DMA, SSRs removed**
- **Memories reduced:**
 - FP64 comparison: 4-bank **TCDM** (64x256), 2-bank **CACHE** (data: 64x128, tag: 41x128)
 - FP32 comparison: 4-bank **TCDM** (32x512), 2-bank **CACHE** (data: 64x128, tag: 27x128)

Comparison

Single-precision support ($T_{clk}=2000ps$):

- 32-bit Snitch + FP32 Fast-FPU*
- 32-bit Snitch + FP32 Tiny-FPU
- 32-bit Snitch

Double-precision support ($T_{clk}=2400ps$):


- 32-bit Snitch + FP64 Fast-FPU*
(merged)
- 32-bit Snitch + FP64 Tiny-FPU
(merged)
- 32-bit Snitch

**Fast-FPU: combinational*

SNITCH


FPU

Snitch



Comparison (Area)

Area Comparison



Next Steps

- Power analysis
- Complete the paper
- Micro-ops